

# Maths for AI

## Lecture 3: Convolution

Matthew Roughan

`<matthew.roughan@adelaide.edu.au>`


School of Mathematical Sciences and AIML,  
University of Adelaide

March 15, 2022

# Convolutions

What is a convolution?

$$f(t) * g(t) = [f * g](t) = \int_{-\infty}^{\infty} f(u) g(t - u) du$$



# Fourier Transform of a Convolution

$$\mathcal{F}\{f_1(t) * f_2(t)\} \rightarrow F_1(s)F_2(s)$$

$$\begin{aligned}\mathcal{F}\{f(t) * g(t)\} &= \mathcal{F}\left\{\int_{-\infty}^{\infty} f(u) g(t-u) du\right\} \\&= \mathcal{F}\left\{\left[\int_{-\infty}^{\infty} f(u) g(t-u) du\right]\right\} \\&= \int_{-\infty}^{\infty} f(u) \int_{-\infty}^{\infty} g(t-u) e^{-i2\pi st} dt du \\&= \int_{-\infty}^{\infty} f(u) G(s) e^{-i2\pi su} du \\&= G(s) \int_{-\infty}^{\infty} f(u) e^{-i2\pi su} du \\&= F(s) G(s)\end{aligned}$$

## Convolution example

Convolution of two rectangular pulses  $r(t)$  where


$$r(t) = u(t + 1/2) - u(t - 1/2), \text{ and } u(t) = \begin{cases} 0, & t < 0 \\ 1, & t \geq 0 \end{cases}$$

$$\begin{aligned} r(t) * r(t) &= \int_{-\infty}^{\infty} r(s) r(t-s) ds \\ &= \begin{cases} 0, & \text{if } t < -1 \\ \int_{-1/2}^{1/2+t} r(s) r(t-s) ds, & \text{if } -1 \leq t \leq 0 \\ \int_{t-1/2}^{1/2} r(s) r(t-s) ds, & \text{if } 0 \leq t \leq 1 \\ 0, & \text{if } t > 1 \end{cases} \end{aligned}$$

## Convolution example

For  $-1 \leq t \leq 0$ , the convolution  $r(t) * r(t)$  is

$$\begin{aligned} r(t) * r(t) &= \int_{-1/2}^{1/2+t} r(s) r(t-s) ds, \\ &= \int_{-1/2}^{1/2+t} 1 ds, \\ &= [s]_{-1/2}^{1/2+t} \\ &= 1/2 + t - (-1/2) \\ &= 1 + t, \end{aligned}$$



Similarly for  $0 \leq t \leq 1$ , we get  $r(t) * r(t) = 1 - t$



## Convolution example

Result is a Triangular pulse


$$r(t) * r(t) = \begin{cases} 0, & \text{if } t < -1 \\ 1 + t, & \text{if } -1 \leq t \leq 0 \\ 1 - t, & \text{if } 0 \leq t \leq 1 \\ 0, & \text{if } t > 1 \end{cases}$$

$\mathcal{F}\{r(t)\} = \text{sinc}(s)$  hence from the convolution theorem


$\mathcal{F}\{r(t) * r(t)\} = \text{sinc}^2(s)$




# Convolution example: interpolation



=



\*



# Convolution example: interpolation

Fourier transformation of a piecewise linear function

$$f(t) = \left[ \sum_{i=1}^n f_i \delta(t - t_i) \right] * r(t) * r(t)$$

is

$$F(s) = \left[ \sum_{i=1}^n f_i e^{-i2\pi s t_i} \right] \text{sinc}^2(s)$$



# Discrete Convolution

$$[x_1 * x_2](n) = \sum_{i=-\infty}^{\infty} x_1(i)x_2(n-i) = \sum_{i=-\infty}^{\infty} x_1(n-i)x_2(i)$$

Now remember the impact of convolutions in DFTs, e.g.

$$\mathcal{F}\{x_1 * x_2\} = X_1(k)X_2(k)$$

where  $\mathcal{F}\{x_1(n)\} = X_1(k)$  and  $\mathcal{F}\{x_2(n)\} = X_2(k)$ .

Discrete convolutions have a special role in signal processing as *Linear Time-Invariant Filters*

# Filters

A filter takes some input  $x(n)$ , and produces an output  $y(n)$ , which has been filtered to extract certain features (e.g. trend, seasonality, ...)



## References:

- Brockwell and Davis, 1996
- Box and Jenkins, 1976
- Anderson and Moore, 1979

# Possible filter properties

- **invertibility:** The mapping  $x(t) \rightarrow y(t)$  must be 1:1, so that each input signal has a unique output signal (don't need to invert all possible outputs).
- **memory:**  $y(t_0)$  depends on  $x(t)$  for  $t \neq t_0$ .
- **causality:**  $y(t_0)$  only depends on  $x(t)$  for  $t \leq t_0$ .
- **stability:** Bounded Input Bounded Output (BIBO). If  $|x(t)| \leq M$  for all  $t$  and some  $M$ , then  $|y(t)| \leq R$  for all  $t$  and some  $R$ .
- **time invariance:** time shift doesn't matter, i.e.  $x(t) \rightarrow y(t)$  implies  $x(t - t_0) \rightarrow y(t - t_0)$ .
- **linearity:** principle of superposition:  $x_i \rightarrow y_i$ ,  $i = 1, 2$  implies that for all  $a_1, a_2 \in \mathbb{R}$ ,  $a_1 x_1 + a_2 x_2 \rightarrow a_1 y_1 + a_2 y_2$ .

# Linear Filters

Response is linear in the input, e.g. given the filter,

$$\mathcal{L}x_1 \rightarrow y_1$$

$$\mathcal{L}x_2 \rightarrow y_2$$

Then

$$\mathcal{L}ax_1 + bx_2 \rightarrow ay_1 + by_2$$

The output of linear filters can be written as a linear combination of the inputs.

$$y(m) = \sum_{i=-\infty}^{\infty} w(m, i)x(m - i)$$

# Linear Time Invariant Filters

- time invariant filters don't change over time, so  $w(m, i) = w(i)$

The output of linear filters can be written as a linear combination of the inputs.

$$y(m) = \sum_{i=-\infty}^{\infty} w(i)x(m-i)$$

Note that this is a discrete convolution!

# Linear Time Invariant Causal Filters

- time invariant filters don't change over time, so  $w(m, i) = w(i)$
- causal filters only depend on the past, so  $w(-i) = 0$ , for  $i > 0$ .

The output of linear filters can be written as a linear combination of the inputs.

$$y(m) = \sum_{i=0}^{\infty} w(i)x(m-i)$$


Note that this is also a discrete convolution!

# Impulse response

Given a filter:



The impulse response is the output of the filter given an impulse as the input.



# Impulse response

For a linear, time-invariant filter  $F$ , the impulse response is

$$I_F(m) = \sum_{i=-\infty}^{\infty} w(i)\delta_{mi} = w(m)$$

where  $\delta_{nk}$  is the Kronecker delta, defined by

$$\delta_{nk} = \begin{cases} 1 & \text{if } n = k \\ 0 & \text{otherwise} \end{cases}$$

So a **linear time-invariant filter** can be completely characterized by its **impulse response**.



# Impulse response

Note that any signal  $x(n)$  can be written as a linear combination of impulses, e.g.

$$x(n) = \sum_{k=-\infty}^{\infty} \delta_{nk} x(k)$$

Given linearity of the filter, the output can be written as the same linear combination of the impulse responses, e.g.

$$\begin{aligned} y(m) &= \sum_{i=-\infty}^{\infty} w(i) \left[ \sum_{k=-\infty}^{\infty} \delta_{m-i,k} x(k) \right] \\ &= \sum_{i=-\infty}^{\infty} w(i) x(m-i) \end{aligned}$$

# Memory

Filters can have finite, or infinite memory

- **FIR:** Finite Impulse Response filters have an impulse response which have a finite number of terms, i.e.  $\exists N$  such that

$$w(n) = 0, \quad \forall |n| > N$$

- **IIR:** Infinite Impulse Response filters have an impulse response with an infinite number of terms.

though for BIBO we require a finite sum, e.g.

$$\sum_{i=-\infty}^{\infty} |w(i)| < \infty$$

# FIR example: Moving Average

(finite) **Moving Average (MA)**

$$y(n) = \sum_{i=-N}^N b(i)x(n-i)$$

typical example, symmetric rectangular windowed MA

$$y(n) = \frac{1}{2N+1} \sum_{i=-N}^N x(n-i)$$

NB: this is a non-causal filter

## FIR example: difference

A **difference operator** (or filter) looks like

$$y(n) = x(n) - x(n - 1)$$

Note this is a special case of the MA above

$$b(0) = 1, b(1) = -1$$

but this terminology is used differently in different fields

- signal processing and stats: MA as defined above
- financial time series: MA  $\Rightarrow$  low pass

NB: this is a causal filter

# Example of IIR filter: EWMA

## Exponentially Weighted Moving Average (EWMA)

$$y(n) = ay(n-1) + (1-a)x(n)$$

alternative IIR representation

$$y(n) = (1-a) \sum_{i=0}^{\infty} a^i x(n-i)$$

gives **exponentially** decreasing weight to historical data

More general case **Autoregressive (AR)** filters

$$y(n) = \sum_{i=1}^p a(i)y(n-i) + b(0)x(n)$$

# Transfer function

- we can represent LTI filter as convolution
- in Fourier domain, convolution becomes a simple product
- LTI filter is completely characterized by FT of its impulse response
- we call the FT of the impulse response the **Transfer function**, e.g.

$$W(k) = DFT(w(n))$$

- The transfer function tells us the impact of the filter on different components of the spectrum of a signal

# Types of filters

- **low pass:** pass low frequencies, stop high frequencies.

these filters act as smoothers of the data.

e.g. EWMA, MA

- **high pass:** pass high frequencies, stop low frequencies.

e.g. differencer – highlights edges

- **band pass:** pass a band of frequencies


- **notch:** exclude a band (sometimes called bandstop)

e.g. remove signal at a particular frequency to prevent feedback (“ringing out”)




Example: MA  $y(n) = \frac{1}{2N+1} \sum_{i=-N}^N x(n-i)$

$f_s = 1000$ ,  $N = 10,000$ , input white noise,  $N = 5$

signal segment



power spectrum




low pass






Example: MA  $y(n) = \frac{1}{2N+1} \sum_{i=-N}^N x(n-i)$

$f_s = 1000$ ,  $N = 10,000$ , input 10 sines evenly spaced freq.

signal segment



power spectrum




low pass




Example: difference  $y(n) = x(n) - x(n-1]$

$f_s = 1000$ ,  $N = 10,000$ , input white noise

signal segment



power spectrum





high pass



Example: difference  $y(n) = x(n) - x(n - 1]$

$f_s = 1000$ ,  $N = 10,000$ , input 10 sines evenly spaced freq.

signal segment



power spectrum




high pass




Example: EWMA  $y(n) = ay(n-1) + (1-a)x(n)$

$f_s = 1000$ ,  $N = 10,000$ , input white noise,  $a = 0.9$

signal segment



power spectrum




low pass




Example: EWMA  $y(n) = ay(n-1) + (1-a)x(n)$

$f_s = 1000$ ,  $N = 10,000$ , input 10 sines evenly spaced freq.

signal segment



power spectrum



low pass

# Do It: What do they sound like?

- Use Colab and Numpy to generate a *white noise* sequence
  - ▶ White means frequency spectrum is flat.
  - ▶ Implied Gaussianity

So this is IID Gaussian.

- Try out some filters and listen to the results
  - ▶ White noise (or static)
  - ▶ Low pass, MA filter, to smooth it
  - ▶ High-pass, difference

We're doing this in Numpy because PyTorch is a little too helpful

We should spend a lot more time talking about noise...

We could do an entire course on filter design

# Filtering in the frequency domain

We could filter thus:

- `fft`
- `filter`
- `ifft`,

but this requires  $O(N \log N)$  operations, which grows non-linearly in  $N$ . For many applications, we can't afford to have filtering operations grow faster than  $O(N)$ , e.g. real-time applications,

- The number of data points will be  $f_s T$
- the time available for computation is  $T$
- time available per data point is  $1/t_s$ , which is constant with respect to  $N$ .

# Convolution in 2D

Convolution generalizes to 2D, e.g., the two-dimensional convolution of continuous functions  $f(x, y) = \delta(x)r(y)$  with  $g(x, y) = r(x)\delta(y)$  is

$$[f * g](x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(x', y') g(x - x', y - y') dx' dy'$$

Likewise, for discrete signals we can extend the idea of a LTI filter (a convolution) to 2D


$$[x * y](n, m) = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} x(k, l) y(n - k, m - l)$$




# Filters in 2D

2D data

signal



power spectrum





# Filters in 2D

Spatial rectangular low-pass

$$\frac{1}{49} \begin{pmatrix} 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 \end{pmatrix}$$


# Filters in 2D

Spatial rectangular low pass, frequency response



# Filters in 2D


Spatial rectangular low pass, frequency response 1D




# Filters in 2D

Spatial rectangular low pass

signal



power spectrum




# Filters in 2D

Approximately Gaussian low-pass

$$\begin{pmatrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 4 & 4 & 4 & 1 \\ 1 & 4 & 12 & 4 & 1 \\ 1 & 4 & 4 & 4 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{pmatrix}$$

# Filters in 2D


Approximately Gaussian low-pass, frequency response




# Filters in 2D

Approximately Gaussian low-pass

signal



power spectrum





# Edge detection

Edge detection is a high pass operation, but there are a number of ways to do this in 2D.

- gradient: look for max and min in gradients in image
  - ▶ Sobel
  - ▶ Prewitt
  - ▶ Roberts
- Laplacian: look for zeros of second derivative
  - ▶ Marrs-Hildreth

# Sobel Edge detection

Look for vertical and horizontal edge separately, using filters

$$\begin{pmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{pmatrix} \quad \text{and} \quad \begin{pmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{pmatrix}$$


The threshold for values above or below  $T$ .

Alternatively, can combine to get edge magnitude and direction by


$$M_{\text{sobel}} = \sqrt{M_{\text{vertical}}^2 + M_{\text{horizontal}}^2}$$
$$\phi_{\text{sobel}} = \tan^{-1}(M_{\text{vertical}}/M_{\text{horizontal}})$$

# Filters in 2D

Sobel edge detection, freq. response  
signal




power spectrum




# Filters in 2D

Sobel edge detection applied to image.

signal




power spectrum




# Filters in 2D

Vertical high pass (Sobel edge detection), freq. response, horizontal



# Filters in 2D


Vertical high pass (Sobel edge detection), freq. response, vertical




# Filters in 2D

Horizontal high pass (Sobel edge detection)

signal




power spectrum




# Filters in 2D

Combined Sobel filters, and thresholded version

signal



thresholded





# Prewitt filters

Look for vertical and horizontal edge separately, using filters

$$\begin{pmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{pmatrix} \quad \text{and} \quad \begin{pmatrix} -1 & -1 & -1 \\ 0 & 0 & 0 \\ 1 & 1 & 1 \end{pmatrix}$$

# Roberts filters

$$\begin{pmatrix} 0 & -1 \\ 1 & 0 \end{pmatrix} \quad \text{and} \quad \begin{pmatrix} 1 & 0 \\ 0 & -1 \end{pmatrix}$$

# Laplacian

Approximation of  $\nabla^2 = \frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2}$

$3 \times 3$

$$\begin{pmatrix} 0 & -1 & 0 \\ -1 & 4 & -1 \\ 0 & -1 & 0 \end{pmatrix}$$


$5 \times 5$

$$\begin{pmatrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 24 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{pmatrix}$$


# Filters in 2D

$5 \times 5$  Laplacian, and its frequency response

signal




power spectrum




# Filters in 2D

$5 \times 5$  Laplacian applied to image

signal




power spectrum



# Embossing an image

Fancy effects from simple filters, e.g. take  $\theta = \pi/6$  and

$$\text{image} = M_{\text{sobel}}^{\text{horizontal}} \cos(\theta) + M_{\text{sobel}}^{\text{vertical}} \sin(\theta)$$



# Embossing an image

